

CME 2001

Data Structures and Algorithms

Zerrin Işık
zerrin@cs.deu.edu.tr

B-Trees

B-Trees

- B-trees are balanced search trees and designed to work well on disks or other direct access secondary storage devices.
- Many database systems use B-trees, or variations to store information.
- The idea of the B-tree also motivates the design of many other disk-based index structures.

B-Trees

- When data volume is large and does not fit in memory, a B-tree- is used.
- The B-tree is always balanced (all leafs appear at the same level)
- Since each disk access exchanges a whole block of information between memory and disk, a node of the B-tree is expanded to hold more than two child pointers, up to the block capacity.
- The B-tree requires that every node (except the root) has to be at least half full.
- An exact match query, insertion, deletion need to access $O(\log_B n)$ nodes, where
 - ***B**: the page capacity in number of child pointers*
 - ***n**: is the number of objects*

Disk –Based Environment

- The computer CPU deals directly with the primary storage (main memory).
- We can access data stored in main memory quickly, but cannot store everything in memory:
 - because memory is expensive.
 - memory is volatile, i.e. if there is a power failure, information stored in memory gets lost.
- The secondary storage stands for magnetic disks. Although it has slower access, it is less expensive and it is non-volatile.

Disk –Based Environment

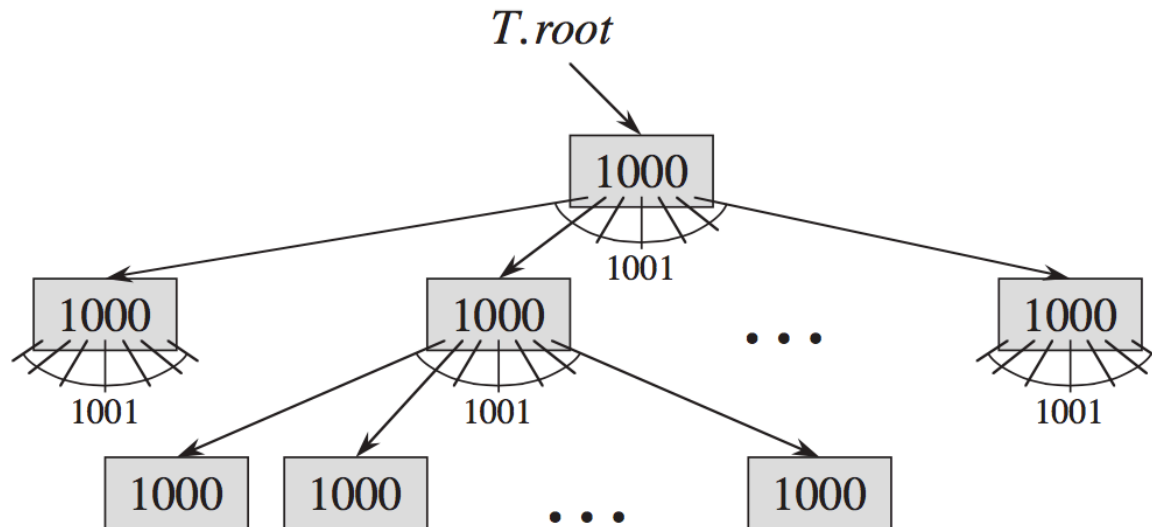
- The CPU does not deal with disk directly, any data has to be read from disk to memory first.
- Data is stored on disk in units called *blocks* or *pages*.
- If a disk page is 8KB (8192 bytes), while a node in the BST is 16 bytes (four integers: key, value, two child pointers)
=> every page is only $8192/16 = 0.2\%$ full.
- To improve space efficiency, we should store multiple tree nodes in one disk page.

Disk –Based Environment

- The running time of a B-tree operations highly depends on the number of *DISK-READ* and *DISK-WRITE* operations.
- ***Branching factor***: total # of children of a parent node
- For a large B-tree stored on a disk, ***branching factors*** are between 50 and 2000, depending on the size of a key relative to the size of a page.
- A large *branching factor* reduces :
 - the height of the tree
 - the number of disk accesses required to find any key

B-tree example

- B-tree :
 - branching factor = 1001
 - height = 2 (that can store over *one billion keys*)
- can keep the root node permanently in main memory,
- can find any key in this tree by making at most only **2** disk accesses.



1 node,
1000 keys

1001 nodes,
1,001,000 keys

1,002,001 nodes,
1,002,001,000 keys

Definition of B-trees

1. Every node x has the following attributes:

- a) $x.n$: the number of keys currently kept in node x
- b) the $x.n$ keys themselves stored in non-decreasing order,
so that $x.key_1 \leq x.key_2 \leq \dots \leq x.key_{x.n}$
- c) $x.leaf$: boolean value, which is *TRUE* if x is a leaf and *FALSE* if x is an internal node.

2. Each internal node x contains $x.n+1$ pointers $x.c_1, x.c_2, \dots, x.c_{n+1}$ to indicate its children. Leaf nodes have no children, no c_i attributes.


3. The keys $x.key_i$ separate the ranges of keys stored in each subtree. If k_i is any key stored in the subtree with root $x.c_i$ then:

$$k_1 \leq x.key_1 \leq k_2 \leq x.key_2 \leq \dots \leq x.key_{x.n} \leq k_{x.n+1}$$

Definition of B-trees


4. All leaves have the same depth, which is the tree's height h .
5. Nodes have lower and upper bounds on the number of keys they can contain. These bounds are represented by a fixed integer $t \geq 2$, called the *minimum degree* of the B-tree:
 - a) Every node (except root) must have at least $t-1$ keys. Every internal node has at least t children. If tree is non-empty, root must have at least one key.
 - b) Every node may contain at most $2t-1$ keys. An internal node may have at most $2t$ children. A node becomes *full*, if it contains exactly $2t-1$ keys.

B-tree example



Insert : 14, 41

B-tree example



B-tree Operations: Search

B-TREE-SEARCH(x, k)

```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

- B-TREE-SEARCH applies a similar method with BST-SEARCH.
- Starts with B-TREE-SEARCH ($T.root, k$)
- If k is found, it returns (x, i) node x and index i such that $x.key_i = k$
- Otherwise returns *NIL*.
- DISK-READ reads the requested object from secondary memory and puts into main memory.

Running time: $O(t.h) = O(t \lg_t n)$

- *while* loop runs t times
- recursion runs at most the depth of the tree = $O(h)$

Create an empty B-tree


B-TREE-CREATE(T)

```
1   $x = \text{ALLOCATE-NODE}()$ 
2   $x.\text{leaf} = \text{TRUE}$ 
3   $x.n = 0$ 
4   $\text{DISK-WRITE}(x)$ 
5   $T.\text{root} = x$ 
```

- ALLOCATE-NODE allocates one disk page to be used as a new node in $O(1)$ time.
- DISK-WRITE saves attributes of new node to the disk.
- Running time: $O(1)$


Insert a key into a B-tree

- To insert new key X , find a suitable leaf node:
 - If leaf node is not full, put X into empty slot (easy case!)
 - If leaf node is full, split leaf node and adjust parents up to root.
- Insert “S” into below B-tree in which $t=4$




Insert a key into a B-tree ...

- If the key should be put into a full root node, which should be divided into two new nodes and new root.
- Insert “H” into the root node of below B-tree in which $t=4$



Insert a key into a B-tree

B-TREE-INSERT(T, k)

```
1   $r = T.root$ 
2  if  $r.n == 2t - 1$   Is Root node is full?
3       $s = \text{ALLOCATE-NODE}()$  If “yes”, then run lines 3-9
4       $T.root = s$ 
5       $s.leaf = \text{FALSE}$ 
6       $s.n = 0$ 
7       $s.c_1 = r$ 
8      B-TREE-SPLIT-CHILD( $s, 1$ )
9      B-TREE-INSERT-NONFULL( $s, k$ )
10 else B-TREE-INSERT-NONFULL( $r, k$ )
```

Running time: $O(t \lg_t n)$

B-TREE-SPLIT-CHILD(x, i)

```
1   $z = \text{ALLOCATE-NODE}()$ 
2   $y = x.c_i$ 
3   $z.\text{leaf} = y.\text{leaf}$ 
4   $z.n = t - 1$ 
5  for  $j = 1$  to  $t - 1$ 
6       $z.\text{key}_j = y.\text{key}_{j+t}$ 
7  if not  $y.\text{leaf}$ 
8      for  $j = 1$  to  $t$ 
9           $z.c_j = y.c_{j+t}$ 
10  $y.n = t - 1$ 
11 for  $j = x.n + 1$  downto  $i + 1$ 
12      $x.c_{j+1} = x.c_j$ 
13  $x.c_{i+1} = z$ 
14 for  $j = x.n$  downto  $i$ 
15      $x.\text{key}_{j+1} = x.\text{key}_j$ 
16  $x.\text{key}_i = y.\text{key}_t$ 
17  $x.n = x.n + 1$ 
18  $\text{DISK-WRITE}(y)$ 
19  $\text{DISK-WRITE}(z)$ 
20  $\text{DISK-WRITE}(x)$ 
```


B-TREE-INSERT-NONFULL(x, k)

```
1   $i = x.n$ 
2  if  $x.\text{leaf}$ 
3      while  $i \geq 1$  and  $k < x.\text{key}_i$ 
4           $x.\text{key}_{i+1} = x.\text{key}_i$ 
5           $i = i - 1$ 
6       $x.\text{key}_{i+1} = k$ 
7       $x.n = x.n + 1$ 
8       $\text{DISK-WRITE}(x)$ 
9  else while  $i \geq 1$  and  $k < x.\text{key}_i$ 
10      $i = i - 1$ 
11      $i = i + 1$ 
12      $\text{DISK-READ}(x.c_i)$ 
13     if  $x.c_i.n == 2t - 1$ 
14          $\text{B-TREE-SPLIT-CHILD}(x, i)$ 
15         if  $k > x.\text{key}_i$ 
16              $i = i + 1$ 
17      $\text{B-TREE-INSERT-NONFULL}(x.c_i, k)$ 
```


Insertion example

min. $t = 3$


(a) initial tree



(b) B inserted




(c) Q inserted




Insertion example

(d) *L* inserted





(e) *F* inserted



Delete a key

There could be several cases while delete a key k from B-tree



Case 1: If k is in node x and x is a leaf



Delete a key ...

Case 2: If k is in node x and x is an internal node:

2a) If the child y that precedes k in node x has at least t keys, then find the predecessor of k in the subtree rooted by y .




Delete a key ...


Case 2b) If child y has fewer than t keys, then, symmetrically, examine the child z that follows k in node x . If z has at least t keys, then find the successor of k in the subtree rooted at z .

Delete a key ...

Case 2c) if both y and z have only $t - 1$ keys, merge k and all of z into y , so that x loses both k and the pointer to z , and y now contains $2t - 1$ keys. Then free z and recursively delete k from y .



(d) G deleted: case 2c




Delete a key ...


Case 3: If k does not present in internal node x , determine the root $x.c_i$ of the appropriate subtree that must contain k , if k is in the tree at all. If $x.c_i$ has only $t-1$ keys, execute step 3a or 3b as necessary to guarantee that we descend to a node containing at least t keys. Then finish by recursing on the appropriate child of x .

- 3a) If $x.c_i$ has only $t-1$ keys but has an immediate sibling with at least t keys, give $x.c_i$ an extra key by moving a key from x down into $x.c_i$, moving a key from $x.c_i$'s immediate left or right sibling up into x , and moving the appropriate child pointer from the sibling into $x.c_i$.
- 3b) If $x.c_i$ and both of $x.c_i$'s immediate siblings have $t-1$ keys, merge $x.c_i$ with one sibling, which involves moving a key from x down into the new merged node to become the median key for that node.


Delete a key ...





(e) D deleted: case 3b



(e') tree shrinks in height



Delete a key ...



Next Week Topics

- Elementary Graph Algorithms (Chapter 22)